


STAT 542 Midterm Project Report – Group 11

I. Problem Statement

As instructed, we will perform a supervised learning task focusing on image recognition. We select the Olivetti faces data set as the target of our analysis. Our objective is to accurately predict the identity of individuals featured in the Olivetti faces data set, which was developed by AT&T Laboratories Cambridge. This data set consists of 400 images, featuring 10 unique images of 40 different individuals. The images were captured at various times between April 1992 and April 1994, under different lighting conditions and facial expressions. However, all of the images were taken against a similar dark background, with subjects positioned front and center.



Each observation within the data set is labeled between 0 and 39, with 4,096 data columns representing every pixel in 64x64 images. Originally, the images were stored as unsigned 8-bit integers in a 256-gray-level format, but they were converted into floating values ranging from 0 to 1, to enhance compatibility with machine learning tasks. Our approach to training involves implementing three algorithms: Logistic Regression, Linear Discriminant Analysis (LDA), and Convolutional Neural Networks (CNN), in conjunction with Principal Component Analysis (PCA) to reduce dimensionality. A detailed explanation of each method can be found in the following section.

II. Methodology

1. Principal Component Analysis (PCA)

Principal component analysis (PCA) is a statistical method used to identify patterns in high-dimensional data by reducing the dimensionality of the data while retaining most of the variability in the data set. PCA works by transforming the original variables into a new set of variables, called principal components, which are linear combinations of the original variables. These principal components are chosen in such a way that they capture the maximum amount of variability in the data, while also being uncorrelated with each other. By reducing the dimensionality of the data, PCA can simplify complex data sets and make it easier to visualize and interpret the data. PCA is widely used in data analysis, machine learning, and pattern recognition applications, and is particularly useful when dealing with large data sets with many variables.

Each Principal Component z_i is a linear combination of the original variables x_1, x_2, \dots, x_m with weights given by each column u_i of matrix U :

$$z_i = u_{1i}x_1 + u_{2i}x_2 + \dots + u_{mi}x_m,$$

while U is the rotation matrix and Z is a version of the data rotated in such a way that the resulting principal components are orthogonal.

2. Logistic Regression Analysis

In logistic regression, the dependent variable (or response variable) is binary, usually represented by 0 and 1. The goal of logistic regression is to model the probability that the dependent variable equals 1 as a function of the independent variables (or predictors), which can be continuous, categorical or a mix of both. The logistic regression model uses a logistic function (also called a sigmoid function) to map the input variables to the output probabilities. The logistic regression model estimates the regression coefficients that best fit the training data by minimizing the log loss (also called cross-entropy loss) between the predicted probabilities and the true labels.

Once the model is trained, it can be used to predict the probability of the dependent variable being 1 for new input values, and the predicted probability can be converted into a binary decision using a threshold value (e.g., 0.5). If the predicted probability is above the threshold, the predicted label is 1, otherwise it is 0. In the case of multi-class classification problems, there are two common ways to extend logistic regression for multi-class classification: One-vs-all and Softmax regression approach.

3. Linear Discriminant Analysis (LDA)

LDA (Linear Discriminant Analysis) is a supervised machine learning algorithm that is commonly used for classification tasks. Like logistic regression, it is also used for binary classification problems, where the target variable takes on one of two possible values. However, LDA can also be extended to multi-class classification problems. The goal of LDA is to find a linear discriminant function that can distinguish between the different classes by projecting the data onto a lower-dimensional subspace.

The LDA algorithm works by first computing the mean and variance of each feature for each class in the training set. It then uses these statistics to estimate the between-class and within-class covariance matrices. The between-class covariance matrix measures the distance between the means of the different classes, while the within-class covariance matrix measures the spread of the data within each class. Next, LDA finds the eigenvectors and eigenvalues of the inverse of the within-class covariance matrix multiplied by the between-class covariance matrix. These eigenvectors are used to create the linear discriminant function that maximizes the separation between the classes. Finally, the LDA algorithm uses the linear discriminant function to classify new instances by projecting them onto the subspace defined by the eigenvectors and applying a decision rule to determine the predicted class label.

4. Convolutional Neural Networks (CNNs)

Convolutional Neural Networks (CNNs) are a specialized form of deep learning algorithm created explicitly for tasks in image processing and computer vision. They incorporate a blend of convolutional layers, which detect significant features within the input image, and pooling layers, which reduce the data's dimensionality. Typically, CNNs leverage activation functions and fully connected layers to carry out the final classification process. Within the convolutional layers, filters or kernels scan the image to pinpoint critical features such as edges, corners, and textures. Subsequently, these features are integrated and processed in successive layers to execute more intricate operations like object recognition and classification. An illustration of CNNs algorithm is included in **(Appendix 1)**. CNNs have demonstrated exceptional performance levels in a broad range of image processing applications, including image classification, object detection, and facial recognition.

III. Results and Comparisons

1. PCA

To perform dimensionality reduction while maintaining ideal accuracy, we included principal components up to a point where 90 percent of the variances could be explained, which is 60 components (**Appendix 2**).

2. Logistics Regression and LDA

	Logistics Regression		Linear Discriminant Analysis	
	Accuracy	Run Time	Accuracy	Run Time
With PCA	0.975	<u>0.093</u>	0.975	0.048
Without PCA	0.975	5.121	<u>0.986</u>	<u>0.491</u>

In terms of accuracy, PCA does not improve the result of logistics regression. However, applying PCA sped up the process by a lot. For LDA, applying PCA barely affected the accuracy while reducing run time. Similarly, applying PCA reduced Logistics Regression running time, as well as slightly lowered its accuracy. Without PCA, LDA gave greater accuracy with a fair run time compared to logistics regression.

3. Convolutional Neural Network

	2 layers with 32, 64 nodes	3 layers with 32, 64, 128 nodes
Batch size = 32	<u>0.9750</u>	0.9625
Batch size = 64	0.9375	0.9250

The parameter tuning process would be specified in the next chapter. The optimal model has an accuracy of 0.975 (underlined). Other settings for the optimal model are as follows:

epoch = 25, a **3x3** kernel size of convolutional layer, **max-pooling** subsampling method (size 2x2), **relu** activation function and **adam** optimizer.

4. Comparison

Best Performance	Logistic Regression	LDA	CNN
Accuracy	0.975	<u>0.986</u>	0.975

While LDA had a slight edge, three models performed well on this dataset. The accuracy is comparable. However, we disregarded comparison of time consumption because CNNs, by its nature, usually take longer time than the others.

IV. Difficulties

1. Loss of Accuracy While Performing PCA

From the scree plot (**Appendix 3**), we found that the ability to explain variances is relatively low starting from the eighth component. Therefore, originally we selected the first seven components to fit the model but the accuracy was bad. Finally, we ended up using 60 principal components as stated in the previous chapter.

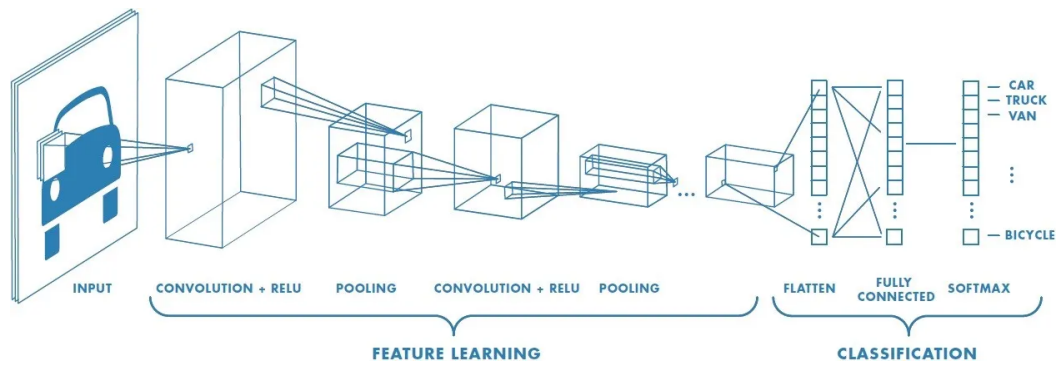
2. Parameter Tuning of CNN

<u>2 hidden layers with 32, 64 nodes</u>	<u>Batch size of 32</u>	<u>25 epochs</u>
3 layers with 32, 64 and 128 nodes	Batch size of 64	60 epochs


The table above consists of the different combinations of parameters that we have gone through. The underlined parameters would produce the optimal model. Some findings during the tuning process include:

- 1) Overfitting: We discovered that the third hidden layer would lead to overfitting (training accuracy = 1). Therefore, the third layer with 128 nodes was dropped.
- 2) The accuracy stopped improving long before epoch reached 60. Based on the graph (**Appendix 4**), we let epoch = 25.


V. Appendices




Appendix 1: Illustration of CNNs algorithm



Appendix 2: Principal Components and Cumulative Variance Explained



Appendix 3: Principal Components and Proportion of Variance Explained



Appendix 4: Number of Epoch and Accuracy

Appendix 5: Code and Result (attached below)



STAT 542: Midterm Project

Face Recognition

- Yu-Ching Liao ycliao3@illinois.edu

Basic Import

Package Import

```
In [1]: # import necessary libraries and modules
from sklearn.datasets import fetch_olivetti_faces
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from keras.wrappers.scikit_learn import KerasClassifier
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
import time
```

Dataset Import

```
In [39]: # load the Olivetti faces dataset
data = fetch_olivetti_faces()
X = data['data']
y = data['target']
```

Prepare train and test set

```
In [40]: # split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran
```

Statistical Learning

Running the Model without PCA

Logistic

```
In [6]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
import seaborn as sns

start = time.time()
# Create a logistic regression model
clf = LogisticRegression()

# Fit the model to the training data
clf.fit(X_train, y_train)

# Use the model to make predictions on the testing data
y_pred = clf.predict(X_test)

# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
end = time.time()


print("Out-sample Accuracy for logistic regression:", accuracy)
print("Time Comsumption:", end - start, "sec.")

cm = confusion_matrix(y_test, y_pred)

fig, ax = plt.subplots(figsize=(15,15))
sns.heatmap(cm, annot=True, fmt='d', ax=ax)

Accuracy for logistic regression: 0.975
Time Comsumption: 6.632861137390137 sec.
```

```
Out[6]: <Axes: >
```

LDA


```
In [55]: #find the best parameters
acc = []
for n in range(1, 40):
    lda = LinearDiscriminantAnalysis(n_components=n)
    lda.fit(X_train, y_train)
    y_pred = lda.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    if n % 5 == 0:
        print("Out-sample Accuracy for LDA for", n, "components:", accuracy)

    acc.append(accuracy)
    accuracy = 0
plt.figure(figsize=(15,8))
plt.title("Accuracies with each number of component")
```

```
plt.plot(range(1,40), acc, 'o-', label = "Accuracy")
plt.ylabel("Accuracies")
plt.xlabel("Number of Components")
plt.legend()
plt.grid(alpha=0.3)
plt.show()
```

```
print("The best number of components:", max(acc))
```

Out-sample Accuracy for LDA for 5 components: 0.9875
 Out-sample Accuracy for LDA for 10 components: 0.9875
 Out-sample Accuracy for LDA for 15 components: 0.9875
 Out-sample Accuracy for LDA for 20 components: 0.9875
 Out-sample Accuracy for LDA for 25 components: 0.9875
 Out-sample Accuracy for LDA for 30 components: 0.9875
 Out-sample Accuracy for LDA for 35 components: 0.9875



The best number of components: 0.9875

So basically all of those number of components have same performance so we are going to randomly pick a number ranged from 1 to 39 and see its performance.

```
In [8]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.metrics import accuracy_score
from mlxtend.plotting import plot_decision_regions

start = time.time()
# Create an instance of LDA
lda = LinearDiscriminantAnalysis(n_components=39)

# Train the LDA model
lda.fit(X_train, y_train)

# Make predictions on the test set
y_pred = lda.predict(X_test)
```


Running the Model with PCA

PCA Visualization

```
In [56]: from sklearn.decomposition import PCA


# load the Olivetti faces dataset
data = fetch_olivetti_faces()
X = data['data']
y = data['target']

# split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

pca = PCA(n_components=20)
X_train_pca = pca.fit_transform(X_train)

# Apply PCA to the testing set
X_test_pca = pca.transform(X_test)

# visualizw results
plt.figure(figsize=(15,8))
plt.plot(range(1, pca.n_components+1), pca.explained_variance_ratio_, 'o-',
plt.title("Accuracies with each number of component")
plt.xlabel('Number of principal components')
plt.ylabel('Proportion of variance explained')
plt.legend()
plt.grid(alpha=0.3)
plt.show()
```



We are going to try the number of components with 3, 5, 7, 10, 14, 20, 30 to see the performance of PCA.

Logistic with PCA

```
In [62]: acc = []
time_comp = []
for n in [3, 5, 7, 10, 14, 20, 30]:
    start = time.time()
    pca = PCA(n_components=n)
    X_train_pca = pca.fit_transform(X_train)
    X_test_pca = pca.transform(X_test)
    clf = LogisticRegression()
    clf.fit(X_train_pca, y_train)
    y_pred = clf.predict(X_test_pca)
    accuracy = accuracy_score(y_test, y_pred)
    end = time.time()
    print("Number of components:", n, ", Out-sample Accuracy for logistic regr")
    print("Time Consumption:", end - start, "sec.\n")
    acc.append(accuracy)
    time_comp.append(end - start)

plt.figure(figsize=(15,8))
plt.plot([3, 5, 7, 10, 14, 20, 30], acc, 'o-', label = 'Accuracy')
plt.ylabel("Accuracies")
plt.xlabel("Number of Components")
plt.title("Accuracies with each number of component")
plt.legend()
plt.grid(alpha=0.3)
plt.show()

plt.figure(figsize=(15,8))
plt.plot([3, 5, 7, 10, 14, 20, 30], time_comp, 'o-', label = 'Time Consumpti')
plt.ylabel("Time Consumption")
plt.xlabel("Number of Components")
plt.title("Time Consumption with each number of component")
plt.legend()
plt.grid(alpha=0.3)
plt.show()
```

Number of components: 3 , Out-sample Accuracy for logistic regression with PCA: 0.3

Time Consumption: 0.5724036693572998 sec.

Number of components: 5 , Out-sample Accuracy for logistic regression with PCA: 0.5125

Time Consumption: 0.5299580097198486 sec.

Number of components: 7 , Out-sample Accuracy for logistic regression with PCA: 0.6375

Time Consumption: 0.3737456798553467 sec.

Number of components: 10 , Out-sample Accuracy for logistic regression with PCA: 0.8

Time Consumption: 0.39716148376464844 sec.

Number of components: 14 , Out-sample Accuracy for logistic regression with PCA: 0.8875


Time Consumption: 0.46981143951416016 sec.


Number of components: 20 , Out-sample Accuracy for logistic regression with PCA: 0.9375

Time Consumption: 0.514904260635376 sec.

Number of components: 30 , Out-sample Accuracy for logistic regression with PCA: 0.9625

Time Consumption: 0.6014773845672607 sec.





It seemed like the differences among time consumptions are not significant so we will try the number of components that have largest accuracy.

```
In [67]: start = time.time()

#Implementing PCA
pca = PCA(n_components=30)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Create a logistic regression model
clf = LogisticRegression()

# Fit the model to the training data
clf.fit(X_train_pca, y_train)

# Use the model to make predictions on the testing data
y_pred = clf.predict(X_test_pca)

# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
end = time.time()


print("Number of components: 30")
print("Out-sample Accuracy for logistic regression with PCA:", accuracy)
print("Time Consumption:", end - start, "sec.")

cm = confusion_matrix(y_test, y_pred)

fig, ax = plt.subplots(figsize=(15,15))
sns.heatmap(cm, annot=True, fmt='d', ax=ax)
```

Number of components: 30
 Out-sample Accuracy for logistic regression with PCA: 0.9625
 Time Consumption: 0.4608933925628662 sec.

Out [67]: <Axes: >



LDA with PCA

```
In [64]: acc = []
time_comp = []
for n in [3, 5, 7, 10, 14, 20, 30]:
    start = time.time()
    pca = PCA(n_components=n)
    X_train_pca = pca.fit_transform(X_train)
    X_test_pca = pca.transform(X_test)
    lda = LinearDiscriminantAnalysis(n_components=n-2)
    lda.fit(X_train_pca, y_train)
    y_pred = lda.predict(X_test_pca)
    accuracy = accuracy_score(y_test, y_pred)
    end = time.time()
    print("Number of components:", n, ", Out-sample Accuracy for LDA with PCA:
```



```

print("Time Consumption:", end - start, "sec.\n")
acc.append(accuracy)
time_comp.append(end - start)

plt.figure(figsize=(15,8))
plt.plot([3, 5, 7, 10, 14, 20, 30], acc, 'o-', label = 'Accuracy')
plt.ylabel("Accuracies")
plt.xlabel("Number of Components")
plt.title("Accuracies with each number of component")
plt.legend()
plt.grid(alpha=0.3)
plt.show()

plt.figure(figsize=(15,8))
plt.plot([3, 5, 7, 10, 14, 20, 30], time_comp, 'o-', label = 'Time Consumpti
plt.ylabel("Time Consumption")
plt.xlabel("Number of Components")
plt.title("Time Consumption with each number of component")
plt.legend()
plt.grid(alpha=0.3)
plt.show()

```

Number of components: 3 , Out-sample Accuracy for LDA with PCA: 0.2875
Time Consumption: 0.07838582992553711 sec.

Number of components: 5 , Out-sample Accuracy for LDA with PCA: 0.5625
Time Consumption: 0.08365988731384277 sec.

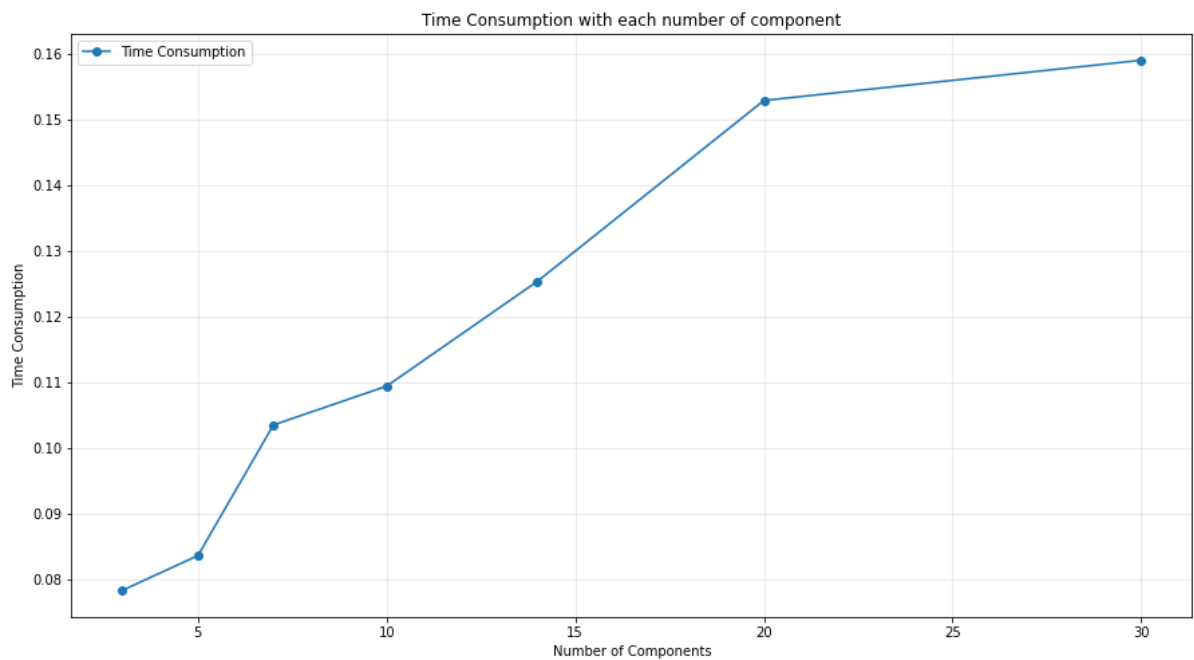
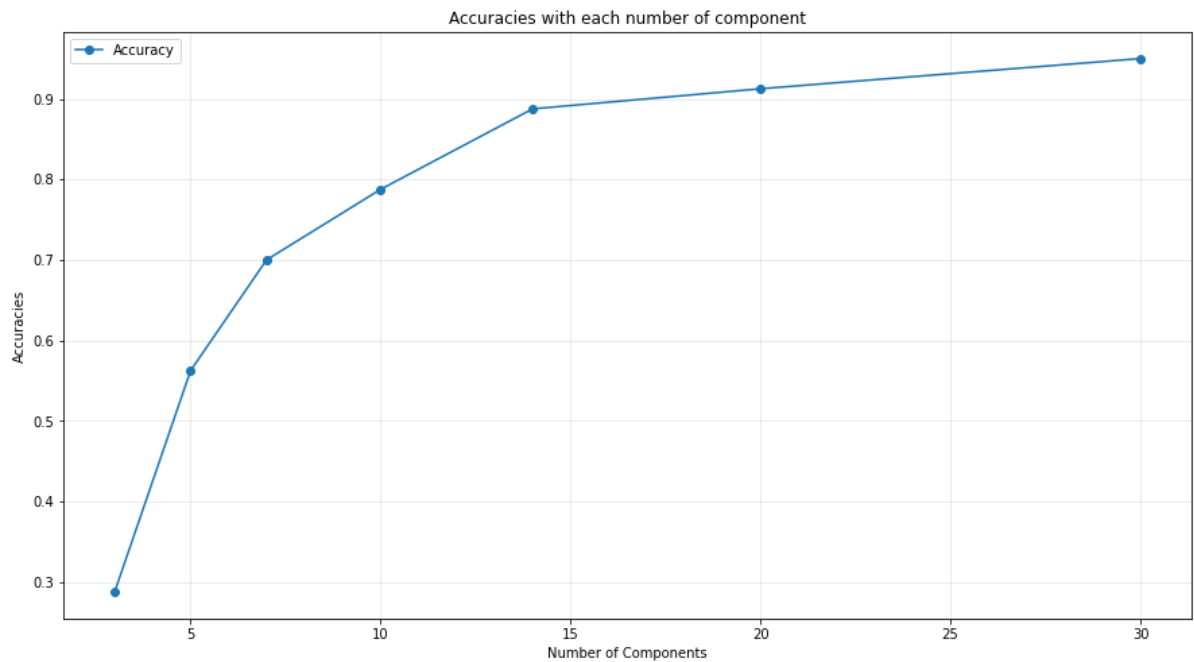
Number of components: 7 , Out-sample Accuracy for LDA with PCA: 0.7
Time Consumption: 0.10351896286010742 sec.

Number of components: 10 , Out-sample Accuracy for LDA with PCA: 0.7875
Time Consumption: 0.10941696166992188 sec.

Number of components: 14 , Out-sample Accuracy for LDA with PCA: 0.8875
Time Consumption: 0.12537312507629395 sec.

Number of components: 20 , Out-sample Accuracy for LDA with PCA: 0.9125
Time Consumption: 0.1528491973876953 sec.

Number of components: 30 , Out-sample Accuracy for LDA with PCA: 0.95
Time Consumption: 0.1589670181274414 sec.



The time consumption will increase nearly twice as the number of components grown from 3 to 30. However, it seemed like the accuracies cease growing since the number of components reached 15. Thus we try number of components with 15 following.

```
In [66]: start = time.time()

#PCA
pca = PCA(n_components=15)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Create an instance of LDA
lda = LinearDiscriminantAnalysis(n_components=5)

# Train the LDA model
```

```

lda.fit(X_train_pca, y_train)

# Make predictions on the test set
y_pred = lda.predict(X_test_pca)

# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)

end = time.time()

# Print the accuracy
print("Number of components: 15")
print("Out-sample Accuracy for LDA:", accuracy)
print("Time Consumption:", end - start, "sec.")

cm = confusion_matrix(y_test, y_pred)

fig, ax = plt.subplots(figsize=(15,15))
sns.heatmap(cm, annot=True, fmt='d', ax=ax)


```

```

Number of components: 15
Out-sample Accuracy for LDA: 0.9
Time Consumption: 0.28126072883605957 sec.

```

Out[66]: <Axes: >



In []:

Deep Learning: CNN

```
In [14]: # Import the necessary libraries
import numpy as np
from sklearn.datasets import fetch_olivetti_faces
from sklearn.model_selection import train_test_split
from tensorflow.keras import layers, models

# Load the Olivetti faces dataset
dataset = fetch_olivetti_faces()
X = dataset.data.reshape(-1, 64, 64, 1) # Reshape the data into images
y = dataset.target
```

```

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

start = time.time()
# Define the CNN model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(40, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(X_train, y_train, epochs=60, batch_size=32, validation_d

# Evaluate the model on the test data
test_loss, test_acc = model.evaluate(X_test, y_test)

end = time.time()
print('Out-sample Test accuracy for CNN:', test_acc)
print("Time Comsumption:", end - start, "sec.")

# Plot the training and validation accuracy history
plt.figure(figsize=(12,8))
plt.plot(history.history['accuracy'], label='Training accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()


```

Epoch 1/60
10/10 [=====] - 3s 228ms/step - loss: 3.7012 - accuracy: 0.0156 - val_loss: 3.6956 - val_accuracy: 0.0000e+00
Epoch 2/60
10/10 [=====] - 2s 206ms/step - loss: 3.6867 - accuracy: 0.0406 - val_loss: 3.7047 - val_accuracy: 0.0000e+00
Epoch 3/60
10/10 [=====] - 2s 204ms/step - loss: 3.6826 - accuracy: 0.0312 - val_loss: 3.7319 - val_accuracy: 0.0000e+00
Epoch 4/60
10/10 [=====] - 3s 328ms/step - loss: 3.6654 - accuracy: 0.0437 - val_loss: 3.7314 - val_accuracy: 0.0000e+00
Epoch 5/60
10/10 [=====] - 2s 212ms/step - loss: 3.6479 - accuracy: 0.0500 - val_loss: 3.7185 - val_accuracy: 0.0000e+00
Epoch 6/60
10/10 [=====] - 2s 210ms/step - loss: 3.5900 - accuracy: 0.0719 - val_loss: 3.6808 - val_accuracy: 0.0250
Epoch 7/60
10/10 [=====] - 2s 209ms/step - loss: 3.3892 - accuracy: 0.1406 - val_loss: 3.5236 - val_accuracy: 0.0500
Epoch 8/60
10/10 [=====] - 2s 210ms/step - loss: 2.9935 - accuracy: 0.2344 - val_loss: 3.0909 - val_accuracy: 0.2375
Epoch 9/60
10/10 [=====] - 2s 222ms/step - loss: 2.2842 - accuracy: 0.4281 - val_loss: 2.5594 - val_accuracy: 0.3625
Epoch 10/60
10/10 [=====] - 4s 423ms/step - loss: 1.5825 - accuracy: 0.5969 - val_loss: 1.6420 - val_accuracy: 0.5750
Epoch 11/60
10/10 [=====] - 4s 370ms/step - loss: 1.1012 - accuracy: 0.6812 - val_loss: 1.0202 - val_accuracy: 0.7375
Epoch 12/60
10/10 [=====] - 4s 379ms/step - loss: 0.6678 - accuracy: 0.8469 - val_loss: 1.0121 - val_accuracy: 0.7375
Epoch 13/60
10/10 [=====] - 4s 420ms/step - loss: 0.4729 - accuracy: 0.8656 - val_loss: 0.9264 - val_accuracy: 0.7375
Epoch 14/60
10/10 [=====] - 4s 355ms/step - loss: 0.3435 - accuracy: 0.8969 - val_loss: 0.4453 - val_accuracy: 0.8625
Epoch 15/60
10/10 [=====] - 2s 225ms/step - loss: 0.2453 - accuracy: 0.9438 - val_loss: 0.4538 - val_accuracy: 0.8250
Epoch 16/60
10/10 [=====] - 2s 211ms/step - loss: 0.1426 - accuracy: 0.9656 - val_loss: 0.4761 - val_accuracy: 0.8875
Epoch 17/60
10/10 [=====] - 2s 215ms/step - loss: 0.1031 - accuracy: 0.9812 - val_loss: 0.3277 - val_accuracy: 0.9125
Epoch 18/60
10/10 [=====] - 3s 337ms/step - loss: 0.0666 - accuracy: 0.9937 - val_loss: 0.3167 - val_accuracy: 0.9250
Epoch 19/60
10/10 [=====] - 2s 207ms/step - loss: 0.0471 - acc

uracy: 0.9937 - val_loss: 0.2262 - val_accuracy: 0.9500
Epoch 20/60
10/10 [=====] - 2s 213ms/step - loss: 0.0283 - accuracy: 1.0000 - val_loss: 0.1460 - val_accuracy: 0.9625
Epoch 21/60
10/10 [=====] - 2s 222ms/step - loss: 0.0189 - accuracy: 1.0000 - val_loss: 0.2014 - val_accuracy: 0.9500
Epoch 22/60
10/10 [=====] - 2s 212ms/step - loss: 0.0112 - accuracy: 1.0000 - val_loss: 0.1818 - val_accuracy: 0.9500
Epoch 23/60
10/10 [=====] - 3s 295ms/step - loss: 0.0091 - accuracy: 1.0000 - val_loss: 0.1586 - val_accuracy: 0.9625
Epoch 24/60
10/10 [=====] - 3s 246ms/step - loss: 0.0072 - accuracy: 1.0000 - val_loss: 0.1492 - val_accuracy: 0.9625
Epoch 25/60
10/10 [=====] - 2s 211ms/step - loss: 0.0068 - accuracy: 1.0000 - val_loss: 0.1343 - val_accuracy: 0.9625
Epoch 26/60
10/10 [=====] - 2s 214ms/step - loss: 0.0051 - accuracy: 1.0000 - val_loss: 0.1554 - val_accuracy: 0.9625
Epoch 27/60
10/10 [=====] - 2s 213ms/step - loss: 0.0046 - accuracy: 1.0000 - val_loss: 0.1454 - val_accuracy: 0.9625
Epoch 28/60
10/10 [=====] - 2s 210ms/step - loss: 0.0041 - accuracy: 1.0000 - val_loss: 0.1333 - val_accuracy: 0.9625
Epoch 29/60
10/10 [=====] - 3s 344ms/step - loss: 0.0036 - accuracy: 1.0000 - val_loss: 0.1405 - val_accuracy: 0.9625
Epoch 30/60
10/10 [=====] - 2s 215ms/step - loss: 0.0033 - accuracy: 1.0000 - val_loss: 0.1398 - val_accuracy: 0.9625
Epoch 31/60
10/10 [=====] - 2s 214ms/step - loss: 0.0030 - accuracy: 1.0000 - val_loss: 0.1377 - val_accuracy: 0.9625
Epoch 32/60
10/10 [=====] - 2s 211ms/step - loss: 0.0028 - accuracy: 1.0000 - val_loss: 0.1403 - val_accuracy: 0.9625
Epoch 33/60
10/10 [=====] - 2s 216ms/step - loss: 0.0027 - accuracy: 1.0000 - val_loss: 0.1415 - val_accuracy: 0.9625
Epoch 34/60
10/10 [=====] - 3s 300ms/step - loss: 0.0025 - accuracy: 1.0000 - val_loss: 0.1385 - val_accuracy: 0.9625
Epoch 35/60
10/10 [=====] - 3s 252ms/step - loss: 0.0023 - accuracy: 1.0000 - val_loss: 0.1375 - val_accuracy: 0.9625
Epoch 36/60
10/10 [=====] - 3s 296ms/step - loss: 0.0022 - accuracy: 1.0000 - val_loss: 0.1405 - val_accuracy: 0.9625
Epoch 37/60
10/10 [=====] - 4s 359ms/step - loss: 0.0021 - accuracy: 1.0000 - val_loss: 0.1403 - val_accuracy: 0.9625
Epoch 38/60

10/10 [=====] - 2s 211ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.1419 - val_accuracy: 0.9625
Epoch 39/60
10/10 [=====] - 3s 337ms/step - loss: 0.0019 - accuracy: 1.0000 - val_loss: 0.1406 - val_accuracy: 0.9625
Epoch 40/60
10/10 [=====] - 2s 211ms/step - loss: 0.0018 - accuracy: 1.0000 - val_loss: 0.1391 - val_accuracy: 0.9625
Epoch 41/60
10/10 [=====] - 2s 210ms/step - loss: 0.0017 - accuracy: 1.0000 - val_loss: 0.1428 - val_accuracy: 0.9625
Epoch 42/60
10/10 [=====] - 2s 213ms/step - loss: 0.0016 - accuracy: 1.0000 - val_loss: 0.1415 - val_accuracy: 0.9625
Epoch 43/60
10/10 [=====] - 2s 212ms/step - loss: 0.0015 - accuracy: 1.0000 - val_loss: 0.1431 - val_accuracy: 0.9625
Epoch 44/60
10/10 [=====] - 3s 312ms/step - loss: 0.0014 - accuracy: 1.0000 - val_loss: 0.1428 - val_accuracy: 0.9625
Epoch 45/60
10/10 [=====] - 2s 231ms/step - loss: 0.0014 - accuracy: 1.0000 - val_loss: 0.1404 - val_accuracy: 0.9625
Epoch 46/60
10/10 [=====] - 2s 212ms/step - loss: 0.0013 - accuracy: 1.0000 - val_loss: 0.1435 - val_accuracy: 0.9625
Epoch 47/60
10/10 [=====] - 2s 213ms/step - loss: 0.0013 - accuracy: 1.0000 - val_loss: 0.1444 - val_accuracy: 0.9625
Epoch 48/60
10/10 [=====] - 2s 213ms/step - loss: 0.0012 - accuracy: 1.0000 - val_loss: 0.1439 - val_accuracy: 0.9625
Epoch 49/60
10/10 [=====] - 2s 218ms/step - loss: 0.0011 - accuracy: 1.0000 - val_loss: 0.1442 - val_accuracy: 0.9625
Epoch 50/60
10/10 [=====] - 3s 333ms/step - loss: 0.0011 - accuracy: 1.0000 - val_loss: 0.1469 - val_accuracy: 0.9625
Epoch 51/60
10/10 [=====] - 2s 215ms/step - loss: 0.0011 - accuracy: 1.0000 - val_loss: 0.1420 - val_accuracy: 0.9625
Epoch 52/60
10/10 [=====] - 2s 211ms/step - loss: 0.0010 - accuracy: 1.0000 - val_loss: 0.1464 - val_accuracy: 0.9625
Epoch 53/60
10/10 [=====] - 2s 213ms/step - loss: 9.7355e-04 - accuracy: 1.0000 - val_loss: 0.1446 - val_accuracy: 0.9625
Epoch 54/60
10/10 [=====] - 2s 213ms/step - loss: 9.3644e-04 - accuracy: 1.0000 - val_loss: 0.1475 - val_accuracy: 0.9625
Epoch 55/60
10/10 [=====] - 3s 321ms/step - loss: 8.9502e-04 - accuracy: 1.0000 - val_loss: 0.1475 - val_accuracy: 0.9625
Epoch 56/60
10/10 [=====] - 2s 238ms/step - loss: 8.6761e-04 - accuracy: 1.0000 - val_loss: 0.1458 - val_accuracy: 0.9625

Epoch 57/60
 10/10 [=====] - 2s 209ms/step - loss: 8.3691e-04 - accuracy: 1.0000 - val_loss: 0.1482 - val_accuracy: 0.9625
 Epoch 58/60
 10/10 [=====] - 2s 209ms/step - loss: 8.0988e-04 - accuracy: 1.0000 - val_loss: 0.1491 - val_accuracy: 0.9625
 Epoch 59/60
 10/10 [=====] - 2s 223ms/step - loss: 7.7744e-04 - accuracy: 1.0000 - val_loss: 0.1458 - val_accuracy: 0.9625
 Epoch 60/60
 10/10 [=====] - 2s 224ms/step - loss: 7.4348e-04 - accuracy: 1.0000 - val_loss: 0.1464 - val_accuracy: 0.9625
 3/3 [=====] - 0s 42ms/step - loss: 0.1464 - accuracy: 0.9625
 Out-sample Test accuracy for CNN: 0.9624999761581421
 Time Consumption: 203.26114439964294 sec.



The accuracies reached its limit after 20 of epoches, which is 96.25% of accuracy.

I ILLINOIS